

Context & Harness Engineering

Lessons from AKD

Nishan Pantha · NASA-IMPACT

2nd ESA-NASA International Workshop on AI Foundation Models for Earth Observation

Day 4 / Track 1 — GeoAI Agent Tutorial · May 2026

The bottleneck has moved

"Context engineering is the delicate art and science of filling the context window with just the right information for the next step."

— *Andrej Karpathy, June 2025*

The **model** isn't the bottleneck anymore.

The **system around the model** is.

PROMPT → CONTEXT

CONTEXT → HARNESS

Prompt engineering → context engineering

Prompt engineering → context engineering


a string becomes a system — what changes when you scale to scientific agents

PROMPT ENGINEERING	CONTEXT ENGINEERING
a clever string	a system that curates
artefact one prompt template; you tweak words.	artefact a directory of files + a tool surface + guardrails.
unit of work a single LLM call.	unit of work a run: write → select → compress → isolate, then call.
scaling "add another paragraph."	scaling retrieve / compact / partition across sub-agents.
failure mode wrong wording → wrong answer.	failure mode poisoning · distraction · confusion · clash · context rot.
evaluation vibes / spot-check.	evaluation benchmark suites · trajectory evals · reproducibility logs.

The four pillars

Context engineering — four pillars

filling the context window with the right information, every step



attributed to LangChain — the canonical decomposition

How contexts fail

Bigger window ≠ better answers

four ways contexts fail in production

01 Poisoning

A hallucination, bad tool output, or stale fact lands in context — then propagates downstream as if it were ground truth.

02 Distraction

Window so full of "relevant" history that the model loses sight of what the user actually asked. Attention budget gone.

03 Confusion

Irrelevant tool outputs, dead references, or stray instructions muddle the next decision. Signal drowned by noise.

04 Clash

Two pieces of context disagree. System prompt says one thing, retrieved doc the other. The model picks — sometimes wrong.


Context rot

Chroma (2025): every frontier model degrades as input grows. Long ≠ thoughtful.
Treat the window as a finite attention budget, not free storage.

refs: Modarressi et al. "NoLiMa" arXiv:2502.05167 (2025) · Chroma "Context Rot" (2025) · dbreunig.com "How Long Contexts Fail" (2025)

Harness engineering


"~98% of Claude Code is deterministic infrastructure." — the agent loop is a while-loop; the harness does the work.



A capable model + a poor harness loses to a weaker model + a great harness. Reliability lives in the surrounding **deterministic** code — tools, loop, guardrails, state.

AKD doesn't invent — it synthesizes

arrows show what borrows from where; right-side boxes with no arrow are AKD-original



composable guardrails — the granite >> risk_agent >> fact_reasoner chain with and · or · fail-fast operators — and CARE are AKD's own.

AKD today — and how we build it

6 domain agents · two closed-loop workflows (CM1, FM Inference) · live across Earth Science · Astrophysics · Planetary Science · flow.akd.odsi.io


Composable guardrails — *input*: Granite Guardian (safety / HAP) · *output*: Risk Agent + Fact Reasoner (risk · factuality + attribution)

→ How do we actually build these agentic workflows?

Context engineering. Harness engineering. **CARE.**

AKD Labs — the 4-stage workbench

labs.akd.odsi.io · SME approval at every gate · publish to the AKD registry



Every published agent passes through these four stages — **SME approval at every gate**, then **publish to the AKD registry**. *Context engineering is built in* — not a step at the end.

Context engineering @ AKD — CARE

CARE = Collaborative Agent Reasoning Engineering

A staged, artefact-driven discipline. Five phases:


#	Phase	Output
1	Scope	Who, what, when, success criteria
2	Key-info elicitation	Domain knowledge interview transcripts
3	Reasoning policy	The loop the agent will run
4	Prompt architecture	The artifact directory
5	Benchmarking	Test suite + rubric

Context engineering — but as a discipline, not an afterthought.

Artifacts as the knowledge layer

Artifact-driven context — AKD & the tutorial agent share this shape


behavior is written, not coded · each file is consumable knowledge for the LLM



Harness engineering @ AKD

Harness engineering @ AKD — the stack


primitives that wrap every published agent — built once, reused everywhere



Composable guardrails — *input*: Granite Guardian (HAP / safety) · *output*: Risk Agent (domain-policy risk) + Fact Reasoner (factuality + attribution) · *dimensions*: Attribution · Factuality · Groundedness.

Multi-agent stance — single thread per agent; multi-agent only at workflow boundaries, with explicit hand-offs and HITL gates.

MCP — pluggable scientific data




Agent code never touches API keys or HTTP clients. Swap an upstream source by changing one env var. Same MCP server feeds multiple agents — `allowed_tools` differs.

End-to-end: staged workflows in production

Two workflows, one staged pattern

each stage carries its own artifact · context · guardrails; explicit hand-offs; HITL gates where needed



INSTANCES LIVE TODAY

- **CM1** — weather computational modeling
- **FM Inference** — EO foundation-model inference

same harness pattern, two scientific workflows — deployed today

Both AKD workflows follow the same **staged pattern** — each stage carries its own *artifact*, *context*, *guardrails*; hand-offs are explicit; approval gates where needed.

Instances live today — **CM1** (weather computational modeling) · **FM Inference** (EO foundation-model inference)

Takeaways — and the tie-back

Three things to take away

1. **Artifact-driven context** — write the agent's knowledge as files, not code.
2. **Deterministic harness** — the loop is small; the surrounding infra is most of the value.
3. **Evaluation closes the loop** — without benchmarking, "context engineering" is just vibes.

Tie-back to today's tutorial

The Prithvi agent you'll run loads from an **artifact** with `scope.md / contexts/ / tools/ / guardrails/ / reasoning.md / output.md / agents.md`.

Same shape as AKD. Same ideas, applied to EO foundation models.